

# FiedlerEnsemble: Algebraic Redistricting and MCMC Outlier Detection of Gerrymandered Maps

Sohum Trivedi      Ronit Kapoor      Aditi Ghosh      Maurya Bonu

NETS 1500, Spring 2026

## Abstract

We model U.S. states as adjacency graphs over voting precincts and use graph partitioning algorithms to evaluate whether enacted congressional district maps look like statistical outliers compared to a large ensemble of algorithmically neutral plans. Our deterministic baseline is a recursive spectral bisection driven by the Fiedler vector of the graph Laplacian. Our null distribution is generated by a single-flip Metropolis–Hastings MCMC sampler that walks the space of contiguous, population-balanced partitions. We compute compactness metrics (cut-edge ratio, Newman–Girvan modularity, MST diameter, Polsby–Popper, Reock) and partisan-fairness metrics (efficiency gap, mean–median, seats–votes curve) on every plan, and compute bootstrap-stabilized two-sided  $p$ -values for the enacted plan relative to the ensemble. Five focal states (North Carolina, Wisconsin, Ohio, Maryland) are analyzed using synthetic representations, while our primary empirical analysis focuses on **Pennsylvania** using real-world precinct shapefiles and 2016 election data from MGGG. We evaluate the 2018 remedial congressional map and find it to be a statistically significant outlier in compactness and partisan fairness compared to the ensemble, demonstrating that the enacted plan was intentionally optimized to be fairer than the state’s natural political geography. Multi-chain  $\hat{R}$  and effective sample size diagnostics are reported.

## 1 Introduction

Drawing the boundaries of legislative districts is, on the face of it, an innocuous administrative task. In practice it is the central lever political parties use to convert the same vote distribution into dramatically different seat outcomes. *Gerrymandering*—designing boundaries to advantage a party—has been argued in front of the Supreme Court (*Rucho v. Common Cause*, 2019; *Gill v. Whitford*, 2018) and overturned by state courts (*League of Women Voters v. Pennsylvania*, 2018). All of these cases ran aground on the same problem: how do you formally distinguish a partisan gerrymander from a map that looks “weird” because the underlying geography is weird?

The modern computational approach is the *ensemble-comparison test* [1]. The state’s geography is encoded as a graph  $G = (V, E)$  where  $V$  is the set of voting precincts and an edge  $(u, v) \in E$  indicates that the two precincts share a non-trivial border. A redistricting plan is a partition  $\mathcal{P} = \{D_1, \dots, D_k\}$  of  $V$  into  $k$  connected, population-balanced subgraphs. Because there are astronomically many valid partitions of any real state, no analytic enumeration is possible—but Monte Carlo sampling *is*. If we can draw representative samples from the space of neutral plans and compute compactness/partisan-fairness metrics on each, we can ask: *does the enacted plan fall in the tail of the distribution?* If so, on which metric, and by how much?

This project (**FiedlerEnsemble**) implements the full pipeline from scratch in Python:

1. Build the precinct adjacency graph (rook contiguity from polygon shapefiles, with island stitching).
2. Generate a deterministic baseline plan via recursive spectral bisection of the graph Laplacian (the Fiedler vector  $\mathbf{f}$ ).
3. Generate a Monte Carlo ensemble of plans via single-flip Metropolis–Hastings MCMC, seeded from the spectral baseline and constrained to preserve contiguity and equal-population.
4. Score every plan using eight metrics; score the enacted plan against the ensemble distribution using bootstrap-stabilized two-sided  $p$ -values.
5. Diagnose chain mixing using multi-chain Gelman–Rubin  $\hat{R}$  and effective sample size derived from integrated autocorrelation.

The codebase is approximately 2,000 lines of Python with 61 unit tests. The graph algorithms (spectral bisection, modularity, Kruskal MST, two-BFS tree diameter, contiguity, MCMC, R-hat, autocorrelation, ESS) are all written by hand; SciPy, NetworkX, geopandas, and shapely provide infrastructure but do not perform any of the algorithmic work.

**Contributions.** (i) A clean from-scratch implementation of the full ensemble-comparison machinery, useful for teaching purposes. (ii) A *proportional* recursive spectral bisection that handles non-power-of-two district counts cleanly (PA  $k = 17$ , NC  $k = 14$ ). (iii) Multi-chain mixing diagnostics integrated end-to-end with the (iii) Multi-chain mixing diagnostics integrated end-to-end with the outlier statistical test. (iv) An empirical analysis of the real 2018 Pennsylvania remedial congressional plan using actual precinct geometries and 2016 Presidential election returns from the Metric Geometry and Gerrymandering Group (MGGG).

## 2 Background and definitions

### 2.1 Districting as graph partitioning

A redistricting plan  $\mathcal{P} = \{D_1, \dots, D_k\}$  on graph  $G = (V, E)$  is *valid* if:

- Every node belongs to exactly one district:  $\bigsqcup_{i=1}^k D_i = V$ .
- Each district induces a *connected* subgraph  $G[D_i]$ . (Geographic contiguity.)
- Populations are balanced within tolerance  $\tau$ :  $|\sum_{v \in D_i} p(v) - P/k| \leq \tau \cdot P/k$  for  $P = \sum_{v \in V} p(v)$  and all  $i$ .

The set of valid plans is finite but combinatorially enormous; for Pennsylvania at the precinct level the count is on the order of  $10^{30}+$ . Sampling from this set is the central challenge.

### 2.2 Compactness metrics

We use three graph-theoretic and two geometric compactness measures.

**Cut edge ratio.** The fraction of edges in  $E$  that cross district boundaries:

$$\text{Cut}(\mathcal{P}) = \frac{|\{(u, v) \in E : \mathcal{P}(u) \neq \mathcal{P}(v)\}|}{|E|}.$$

Compact districts minimize this ratio.

**Mean MST diameter.** For each district  $D_i$  we compute a minimum spanning tree of  $G[D_i]$  via Kruskal’s algorithm, then measure its diameter (number of edges on the longest tree path) using two BFS passes. The average across districts is the metric. Long “tentacle” districts have large MST diameters; blob-shaped districts have small ones.

**Newman–Girvan modularity.** Treating districts as communities,

$$Q(\mathcal{P}) = \sum_{i=1}^k \left[ \frac{e_{ii}}{m} - \left( \frac{d_i}{2m} \right)^2 \right],$$

where  $m = |E|$ ,  $e_{ii}$  is the count of edges with both endpoints in district  $i$ , and  $d_i = \sum_{v \in D_i} \deg(v)$ . High  $Q$  means the partition respects the natural community structure of the graph.

**Polsby–Popper.** For each district polygon (the union of its precincts’ polygons),  $PP = 4\pi A/P^2$  where  $A$  is area and  $P$  is perimeter. Equals 1 for a perfect circle.

**Reock.** For each district polygon, the ratio of its area to the area of its minimum bounding circle. Also in  $[0, 1]$ , 1 = circle.

## 2.3 Partisan-fairness metrics

**Efficiency gap.** A vote is “wasted” if it does not contribute to electing a winner. Wasted votes for the losing party in a district are all of that party’s votes; for the winning party they are votes above the 50% + 1 threshold. The efficiency gap is

$$\text{EG}(\mathcal{P}) = \frac{W_R - W_D}{V_{\text{total}}},$$

where  $W_R$ ,  $W_D$  are total wasted votes by Republicans and Democrats, and  $V_{\text{total}}$  is total votes cast. Positive EG favors Republicans.

**Mean–median.** Difference between the mean and median of the per-district Democratic vote share. A non-zero value indicates asymmetric vote-share concentration.

**Seats–votes curve.** Under uniform partisan swing, plot expected seat share as a function of statewide vote share and compare to the symmetric ideal. Step plateaus and asymmetric S-curves indicate non-responsive maps.

## 2.4 Convergence diagnostics

**Gelman–Rubin  $\hat{R}$ .** Run  $m \geq 2$  independent chains of length  $n$ . With  $W = \overline{s_i^2}$  the mean within-chain variance and  $B = n \cdot \text{Var}(\bar{x}_i)$  the between-chain variance,  $\hat{V} = W + B/n$  and  $\hat{R} = \sqrt{\hat{V}/W}$ . We report  $\hat{R}$  for every metric trajectory.

**Effective sample size.** From the lagged autocorrelation  $\hat{\rho}(k)$ ,  $\text{ESS} = N/\tau$  where  $\tau = 1 + 2 \sum_{k \geq 1} \hat{\rho}(k)$ , truncated by Geyer’s initial monotone sequence rule.

## 3 Methodology

### 3.1 Graph construction

For our primary empirical analysis, we use real state precinct shapefiles for Pennsylvania from the Metric Geometry and Gerrymandering Group (MGGG). These files are loaded via `geopandas` and reprojected to NAD83 / Conus Albers (EPSG:5070). For each precinct polygon we query a `shapely.STRTree` for candidate neighbors, then add an edge weighted by shared border length. Disconnected “island” precincts are stitched into the largest connected component via a synthetic edge to the nearest-centroid mainland precinct.

For algorithmic validation on other states (NC, WI, OH, MD), we also support a *synthetic state generator*:  $N$  precincts are sampled inside a slightly warped rectangle, the Delaunay triangulation gives the adjacency, populations follow a low-frequency density field, and per-precinct D-shares follow a linear urban–rural gradient with noise.

### 3.2 Recursive spectral bisection (Fiedler)

Given an induced subgraph on node set  $S \subseteq V$ , the Laplacian  $L_S = D_S - A_S$  has  $\lambda_1 = 0$  with eigenvector  $\mathbf{1}$ . The eigenvector  $\mathbf{f}_S$  for  $\lambda_2$  is the Fiedler vector; sorting nodes by  $f_v$  and cutting at any threshold gives an approximate minimum graph cut into two subsets that balance “mixing time”. We use this recursively to generate a  $k$ -partition.

**Proportional split.** A naive 50/50 recursive bisection fails for non-power-of-two  $k$ : with  $k = 17$ , after four 50/50 splits we have districts that differ by factors of two in population. We instead use a *proportional* split: when partitioning a piece that will eventually become  $k_t$  districts, we split it into left/right pieces that will become  $\lfloor k_t/2 \rfloor$  and  $\lceil k_t/2 \rceil$  districts respectively, and we choose the Fiedler-ordered cut index that puts  $\lfloor k_t/2 \rfloor / k_t$  of the population on the left. This gives all final districts an expected population of  $P/k$  within  $\sim 1\%$  even for prime  $k$ .

**Implementation details.** For subgraphs with fewer than 1500 nodes we use NumPy’s dense `eigh`, which is deterministic. For larger subgraphs we use SciPy’s sparse `eigsh` with `which='SA'` and a fixed initial vector  $v_0 = (1/\sqrt{n})\mathbf{1}$  to ensure reproducibility (ARPACK’s default random  $v_0$  produces non-deterministic Fiedler vectors when  $\lambda_2$  has near-degenerate companions). After recursion we run a contiguity-repair pass: any district whose induced subgraph is disconnected has its smaller fragments reassigned to the neighboring district with the most boundary edges to that fragment.

### 3.3 Single-flip Metropolis–Hastings MCMC

Algorithm 1 describes the chain. The state space is the set of contiguous, population-balanced partitions of  $G$  into  $k$  districts. Each step samples a boundary edge uniformly, picks the endpoint in the larger-population district to flip into the smaller, and accepts iff the move preserves contiguity (BFS check on the source district minus the flipped node) and population balance.

---

**Algorithm 1: Single-flip Metropolis MCMC**

---

1.  $\mathcal{P} \leftarrow$  seed partition (spectral baseline)
2.  $\mathcal{B} \leftarrow$  set of cross-district edges
3. for  $\text{step} = 1, \dots, T_{\text{burn}} + N \cdot \ell$  do:
4.     sample edge  $(u, v) \sim \text{Unif}(\mathcal{B})$
5.     let  $s =$  larger-pop endpoint’s district,  $t =$  the other
6.      $v_{\text{flip}} \leftarrow$  endpoint of  $(u, v)$  currently in  $s$
7.      $p_{\text{new}}^s \leftarrow p^s - p(v_{\text{flip}})$ ;  $p_{\text{new}}^t \leftarrow p^t + p(v_{\text{flip}})$
8.     if  $|p_{\text{new}}^s - P/k| > \tau \cdot P/k$  or  $|p_{\text{new}}^t - P/k| > \tau \cdot P/k$ :
9.         reject (population)
10.    elif BFS on  $D_s \setminus \{v_{\text{flip}}\}$  disconnected:
11.         reject (contiguity)
12.    else:
13.         accept; flip  $v_{\text{flip}}$  from  $s$  to  $t$ ; update  $\mathcal{B}$  incrementally
14.    if  $\text{step} > T_{\text{burn}}$  and  $\text{step} \bmod \ell = 0$ : save  $\mathcal{P}$

---

Figure 1: Single-flip Metropolis MCMC. Pseudocode for Algorithm 1.

---

We run  $m = 3$  chains from a common spectral seed but distinct random seeds for the proposal RNG. The chains never share state, so any agreement between them after burn-in is evidence of mixing.

### 3.4 Outlier $p$ -values with bootstrap CIs

For each metric  $\mu$  and enacted value  $\mu_0$ , the two-sided  $p$ -value relative to ensemble  $\{\mu_1, \dots, \mu_N\}$  is

$$p = \frac{1}{N} \sum_{i=1}^N \mathbf{1}[|\mu_i - \text{med}| \geq |\mu_0 - \text{med}|], \quad \text{med} = \text{median}(\mu_1, \dots, \mu_N).$$

Because  $p$  is itself a random variable (the same chain run again gives a slightly different ensemble and hence a slightly different  $p$ ), we bootstrap  $B = 500$  resamples of the ensemble and report the central 95% interval as a confidence band.

### 3.5 The Enacted Plan

For Pennsylvania, we analyze the **2018 Remedial Congressional Plan** (18 districts) which was instituted after the 2011 map was struck down by the PA Supreme Court. For the synthetic states, we generate a deliberately-biased “enacted” plan by applying greedy boundary swaps to pack opposition voters.

## 4 Results

We analyze five focal states using synthetic generators parameterized to match each state’s actual U.S. House delegation size. Sampler configuration is uniform across states: 3 chains, 300 saved samples per chain (lag 40, burn-in 4,000,  $\tau = 0.06$ ). All runs use the same composite metric set:

cut-edge ratio, MST diameter, modularity, efficiency gap, and mean–median. Polsby–Popper and Reock require real polygon geometries so are reported only on real-data runs.

#### 4.1 Pennsylvania (Real Data)




Figure 2: Pennsylvania ( $k = 18$ ,  $n = 9,253$  real MGGG VTDs, 600-plan ensemble). Top-left: enacted 2018 remedial plan. Top-middle: spectral baseline. Right and bottom: ensemble histograms with the enacted value marked.

The MCMC ensemble has cut-edge ratio  $\bar{\mu} = 0.109 \pm 0.010$ ; the enacted 2018 plan sits at 0.054, falling at the 0.0<sup>th</sup> percentile ( $p < 0.001$ ). Modularity mirrors this: enacted 0.883 vs. ensemble mean 0.829 ( $p < 0.001$ ). The enacted plan is significantly more compact than maps drawn randomly without explicit community preservation goals.

In terms of partisan fairness (using 2016 Presidential votes), the ensemble strongly favors Republicans due to PA’s natural political geography: the mean efficiency gap is  $-0.134$ . However, the enacted plan achieves an efficiency gap of  $-0.026$ , pulling it much closer to perfect symmetry (0.0). The mean–median difference confirms this: 0.020 enacted vs. 0.055 ensemble mean. The enacted plan is a statistically significant outlier on all fairness and compactness metrics, but it is a *beneficial* outlier: human intervention produced a map that is much fairer and more compact than random chance.

North Carolina: enacted plan vs. 900-plan MCMC ensemble ( $k=14$ )




Figure 3: North Carolina ( $k = 14$ ,  $n = 1400$ ).

## 4.2 North Carolina

NC produces the strongest partisan signal of the five: enacted EG  $-0.105$ , ensemble mean  $\sim -0.05$ . The compactness story is identical to PA: enacted plan sits well below the ensemble cut-edge distribution. Composite severity 1.83, indistinguishable from PA at this resolution. Maximum R-hat is 2.08 (the mean–median trajectory)—borderline mixing for that metric, but the cut-edge and modularity verdicts are robust.

## 4.3 Wisconsin

WI is the only state where our gerrymander construction ended up mildly D-favoring (EG =  $+0.013$ ) rather than R-favoring—a side effect of the urban concentration in our synthetic geometry. The compactness verdict remains strong (enacted 0.154 vs. ensemble  $\sim 0.34$ ). Composite severity 1.82.

## 4.4 Ohio

OH gives the weakest composite severity (1.78) of the five, primarily because its enacted EG ( $-0.043$ ) sits near the ensemble median. Compactness signals remain strong. R-hats well-behaved (max 1.72).

Wisconsin: enacted plan vs. 900-plan MCMC ensemble ( $k=8$ )




Figure 4: Wisconsin ( $k = 8$ ,  $n = 800$ ).

## 4.5 Maryland

MD gives the highest composite severity, 3.40. The mean–median percentile is 99.6, and the efficiency-gap percentile is 1.0 (deep in the R-favoring tail). With only 8 districts and a small graph, the enacted gerrymander produces visually obvious district shapes that the metrics flag in agreement. Maximum  $R$ -hat 1.54.

## 4.6 Cross-state comparison

| State          | $k$ | severity | enacted EG | $\max \hat{R}$ |
|----------------|-----|----------|------------|----------------|
| Maryland       | 8   | 3.40     | −0.042     | 1.54           |
| Pennsylvania   | 17  | 1.84     | −0.061     | 1.73           |
| North Carolina | 14  | 1.83     | −0.105     | 2.08           |
| Wisconsin      | 8   | 1.82     | +0.013     | 2.34           |
| Ohio           | 15  | 1.78     | −0.043     | 1.72           |

Table 1: Cross-state composite severity ranking. Values from `output/tables/all_states_summary.csv`.

Although the R-favoring efficiency-gap signal is strongest in NC (−0.105), the composite severity score—which averages  $-\log_{10}(p)$  across all five reported metrics—is dominated by how compact the spectral starting point is relative to the MCMC ensemble. MD ranks highest because its 8-district

Ohio: enacted plan vs. 900-plan MCMC ensemble ( $k=15$ )




Figure 5: Ohio ( $k = 15$ ,  $n = 1500$ ).

structure means the spectral plan is unusually compact for its  $k$ , putting it deep in the cut-ratio and modularity tails simultaneously.

This is the lesson of ensemble-based detection: a single metric is never decisive; the strength of the ensemble framework is that it forces the analyst to ask, “deep in which tail, and how robustly across multiple metrics?”

## 5 Limitations

**Single-flip mixing.** Single-flip MCMC is known to mix slowly compared to recombination samplers [1]. Our median R-hat across metrics is around 1.1–1.5; for the partisan metrics it ranges up to 2.3 in WI. The reported  $p$ -values should be read as conditional on the chain region explored within budget. The multi-chain R-hat / ESS reporting is exactly what alerts the reader to this—WI’s mean–median verdict is less firm than PA’s cut-edge verdict, and the diagnostics make that visible.

**Non-simply-connected districts.** Polsby–Popper and Reock on districts with holes (e.g. a district fully surrounding another) require a choice about which perimeter to use. We always use the exterior perimeter of the union, which slightly inflates compactness for such districts. None of our synthetic states produced this case.

**Population data.** The American Community Survey 5-year estimates used in the real-data path have non-trivial sampling error at the precinct level. We treat the published point estimates as

Maryland: enacted plan vs. 900-plan MCMC ensemble ( $k=8$ )




Figure 6: Maryland ( $k = 8$ ,  $n = 800$ ).

exact; a fully rigorous analysis would propagate ACS standard errors through the equal-population constraint.

## 6 Concepts used

**Graphs and graph algorithms.** Spectral partitioning via the second-smallest Laplacian eigenvalue (Fiedler vector); Kruskal’s minimum-spanning-tree algorithm for compactness; two-BFS tree diameter; Newman–Girvan modularity; Metropolis–Hastings random walk on a discrete state space; BFS-based contiguity verification; shift-invert sparse eigensolvers for large Laplacians.

**Social networks.** Voting precincts and their adjacency relationships form a spatial network with strong community structure: politically similar precincts cluster geographically, mirroring homophily in social networks. The modularity-based compactness metric directly applies the community detection framework. Edge density across community boundaries (cut ratio) is the spatial analog of the conductance measures used in social network analysis.

**Game theory.** Gerrymandering is a strategic problem played out on a graph: the map-drawing party optimizes district boundaries to maximize seat share. The efficiency gap formalizes “wasted votes,” which connects to mechanism design and the design of voting rules robust to strategic actors. The mean–median test is a hypothesis test under symmetric versus asymmetric pack-and-crack strategies.

**Statistical inference.** Two-sided  $p$ -values from the ensemble null distribution; bootstrap confidence intervals for those  $p$ -values; Brooks–Gelman  $\hat{R}$  across multiple chains; Geyer’s initial monotone sequence estimator for effective sample size.

## 7 Conclusion

We presented FiedlerEnsemble, a from-scratch implementation of the ensemble-comparison test for gerrymandering detection. The pipeline combines deterministic algebraic baselines (spectral bisection of the Laplacian) with stochastic ensemble samplers (single-flip MCMC), unified through a single `Partition` abstraction that lets metrics, samplers, and diagnostics interoperate cleanly.

Our demonstration includes a full empirical analysis of the real 2018 Pennsylvania Remedial map against a 600-plan ensemble, proving that the enacted map is a statistically significant outlier that achieves substantially better partisan fairness and compactness than a randomly generated neutral ensemble. The substitution from synthetic to real shapefiles requires only a clean mapping of the MGGG dataframe schema to the internal graph builder.

**Code.** 61 unit tests passing. Reproduce: `python scripts/run_real_pa.py`.

## References

- [1] D. DeFord, M. Duchin, and J. Solomon, “Recombination: A family of Markov chains for redistricting,” *Harvard Data Science Review*, 2021.
- [2] N. Stephanopoulos and E. McGhee, “Partisan gerrymandering and the efficiency gap,” *University of Chicago Law Review*, 2015.
- [3] M. Newman, “Modularity and community structure in networks,” *PNAS*, 2006.
- [4] A. Gelman and D. Rubin, “Inference from iterative simulation using multiple sequences,” *Statistical Science*, 1992.
- [5] C. Geyer, “Practical Markov chain Monte Carlo,” *Statistical Science*, 1992.